Chandra Akash Kiran

Data Scientist practitioners with hands-on experience in Python, SQL, Pandas, NumPy, and scikit-learn. Builds machine learning models and AI apps using LLMs and RAG. Built a fraud detection model with 97.81% ROC-AUC and 82.54% F1-score, and a full-stack app that answers questions from PDF files. Analyzed 100K+ records.

c9a9kiran@gmail.com 8602201879 India

Education

MTECH — MATS University, Raipur, Chhattisgarh
Computer Science — 2013 - 2016

BE — Chhattisgarh Swami Vivekananda University, Bhilai, Chhattisgarh
Computer Science — 2010 - 2013

Experience

GrowthPaay — Indore, India
Data Analyst Intern — January 2023 - July 2023

- Collected and cleaned 100K+ records using Python, BeautifulSoup, Selenium, REST APIs, Pandas, and NumPy for reports and analysis.
- Wrote SQL queries and automated reports, cutting report time by 40%, and shared EDA findings on trends with stakeholders.

Natrinai Technology Pvt Ltd — Chennai, India
System Engineer — May 2019 - June 2020

- Wrote Python scripts to extract, process, check, and report data, cutting manual work by 50%.
- Built test automation with Python, Selenium, and Robot Framework; analyzed test results, logs, and defect data to support root-cause analysis.

Planning and Statistical Department — Dhamtari, India
Data Operator — October 2014 - March 2016

- Managed, cleaned, and checked large government datasets in Excel, and prepared Pivot Table reports to support decisions based on statistics.

Projects

Enterprise GPT - AI Knowledge Assistant — September 2026 - October 2026

- Finding answers inside long PDF files is slow. Built and deployed a full-stack app that answers questions from uploaded PDFs using RAG.
- Built the pipeline: PDF upload, text extraction, chunking, Gemini embeddings, ChromaDB search, and source-based answers using LangChain agents with Gemini and OpenAI.
- Added secure login, chat history, and per-user document access with Supabase and PostgreSQL, deployed Next.js on Vercel and FastAPI on Render.

Credit Card Fraud Detection — May 2026 - June 2026

- Fraud is rare in credit card data, so most models miss it. Built and deployed a fraud detection model on this imbalanced data.
- Cleaned the data, balanced the classes with SMOTE, compared Logistic Regression with Random Forest, and tuned the threshold to 0.68.
- Random Forest reached 94.55% precision, 73.24% recall, 82.54% F1-score, and 97.81% ROC-AUC, shown in a Streamlit dashboard.

Zomato SQL Analytics - Database and Business Intelligence — April 2026

- Order data needs a clear structure. Designed a relational database with primary and foreign keys for customers, restaurants, menu items, orders, payments, deliveries, and reviews.
- Wrote SQL with JOINs, aggregations, subqueries, CTEs, and window functions to study revenue, customer behavior, restaurant performance, and delivery speed.
- Built and deployed a Streamlit app on Neon PostgreSQL that shows KPIs, trends, and business insights.

Profiles

GitHub
LinkedIn
Portfolio

Skills

Programming and Data
Python, Pandas, NumPy, SQL, Excel

Machine Learning
scikit-learn, Logistic Regression, Random Forest, SMOTE, Classification, Feature Engineering, Model Evaluation

Data Science
Data Cleaning, Data Preprocessing, EDA, Statistical Analysis, Data Validation, Imbalanced Data, Precision, Recall, F1-score, ROC-AUC

GenAI
LLMs, RAG, LangChain, Google Gemini, OpenAI, Prompt Engineering, Embeddings, ChromaDB, Vector Search, AI Agents

Visualization
Matplotlib, Plotly, Power BI, Tableau, Streamlit, Dashboards

Backend and Data Engineering
ETL, REST APIs, FastAPI, Supabase, AWS S3, Web Scraping, BeautifulSoup, Selenium

Tools and Cloud
Git, GitHub, Vercel, Render

Certifications

IBM Data Science Professional Certificate
December 2021
IBM

Data Analytics Training and Certification
February 2023
Masai School